

AI Model Landscape Analysis

A comprehensive review of 10 major AI model providers and key model releases within the last 12 months.

1. OpenAI — GPT-5.2 Pro

Release Date:	December 11, 2025
Knowledge Cut-off:	August 2025
Cost of Tokens:	\$15.00 / M input \$60.00 / M output (Est. Pro Tier)
Context Window:	400K tokens (128K output tokens)
Availability:	Public / Paid (via ChatGPT plans & Developer API)

When to use this model: *Advanced multi-step project execution, complex data synthesis, and maximum accuracy tasks where reasoning depth out-values raw generation speed.*

2. Anthropic — Claude Opus 4.7

Release Date:	April 16, 2026
Knowledge Cut-off:	Late 2025
Cost of Tokens:	\$15.00 / M input \$75.00 / M output
Context Window:	200K tokens
Availability:	Public / Paid Tier (Claude.ai & API)

When to use this model: *Extreme, graduate-level conceptual reasoning, comprehensive long-form writing, and highly precise academic or financial research.*

3. Google DeepMind — Gemini 2.0 Pro

Release Date:	February 5, 2025
Knowledge Cut-off:	August 2024
Cost of Tokens:	\$7.00 / M input \$21.00 / M output (Pay-as-you-go standard)
Context Window:	2 Million tokens
Availability:	Public (Google AI Studio, Vertex AI, & Gemini Advanced)

When to use this model: *Analyzing massive payloads of information like entire complex codebases, hours of video content, or massive enterprise documentation in a single prompt.*

4. Meta AI — Llama 3 (405B Variant)

Release Date:	July 23, 2024
Knowledge Cut-off:	December 2023
Cost of Tokens:	Free to self-host ~ \$0.53 / M input on managed APIs
Context Window:	128K tokens
Availability:	Public / Open-Weights

When to use this model: *Enterprise-grade private setups, synthetic data generation, and highly-tailored fine-tuning on proprietary local hardware.*

5. Mistral AI — Pixtral Large

Release Date:	Late 2024 / Updated Early 2025
Knowledge Cut-off:	Mid 2024
Cost of Tokens:	\$2.00 / M input \$6.00 / M output

Context Window:	128K tokens
Availability:	Public / Open-Weights & Managed API

When to use this model: *High-fidelity, multi-lingual European operations requiring native processing of dense document imagery, charts, and diagrams combined with strong text reasoning.*

6. Microsoft — Phi-4

Release Date:	January 2025
Knowledge Cut-off:	Late 2024
Cost of Tokens:	Free to self-host ~ \$0.10 / M tokens on Azure
Context Window:	32K tokens
Availability:	Public / Open-Weights

When to use this model: *High-speed edge computing, mobile-device deployment, and robust logical reasoning tasks natively contained on local architecture.*

7. Cohere — Command R+ (2025 Edition)

Release Date:	Early 2025
Knowledge Cut-off:	Late 2024
Cost of Tokens:	\$2.50 / M input \$10.00 / M output
Context Window:	128K tokens
Availability:	Public / Commercial API

When to use this model: *Production-grade Retrieval-Augmented Generation (RAG) tasks across complex, multi-lingual enterprise workflows and tool-use automation.*

8. DeepSeek — DeepSeek-V3

Release Date:	December 2024
Knowledge Cut-off:	July 2024
Cost of Tokens:	\$0.14 / M input \$0.28 / M output (Ultra-low cost)
Context Window:	128K tokens
Availability:	Public API / Open-Weights

When to use this model: *High-volume text transformations, code generation pipelines, and analytical processing on tightly constrained technical budgets.*

9. X.AI — Grok 2

Release Date:	August 2024
Knowledge Cut-off:	Mid 2024
Cost of Tokens:	\$2.00 / M input \$10.00 / M output (via API)
Context Window:	128K tokens
Availability:	Public (X Premium Access & Developer API)

When to use this model: *Real-time information routing requiring live synthesis of social trends, public news feeds, and unrestricted creative generation.*

10. Alibaba Cloud — Qwen2.5-72B-Instruct

Release Date:	September 2024
---------------	----------------

Knowledge Cut-off:	August 2024
Cost of Tokens:	Free to self-host ~ \$0.40 / M tokens on public hosts
Context Window:	128K tokens
Availability:	Public / Open-Weights

When to use this model: *High-performance coding tasks, mathematical troubleshooting, and top-tier performance for English-Chinese multi-lingual localization projects.*